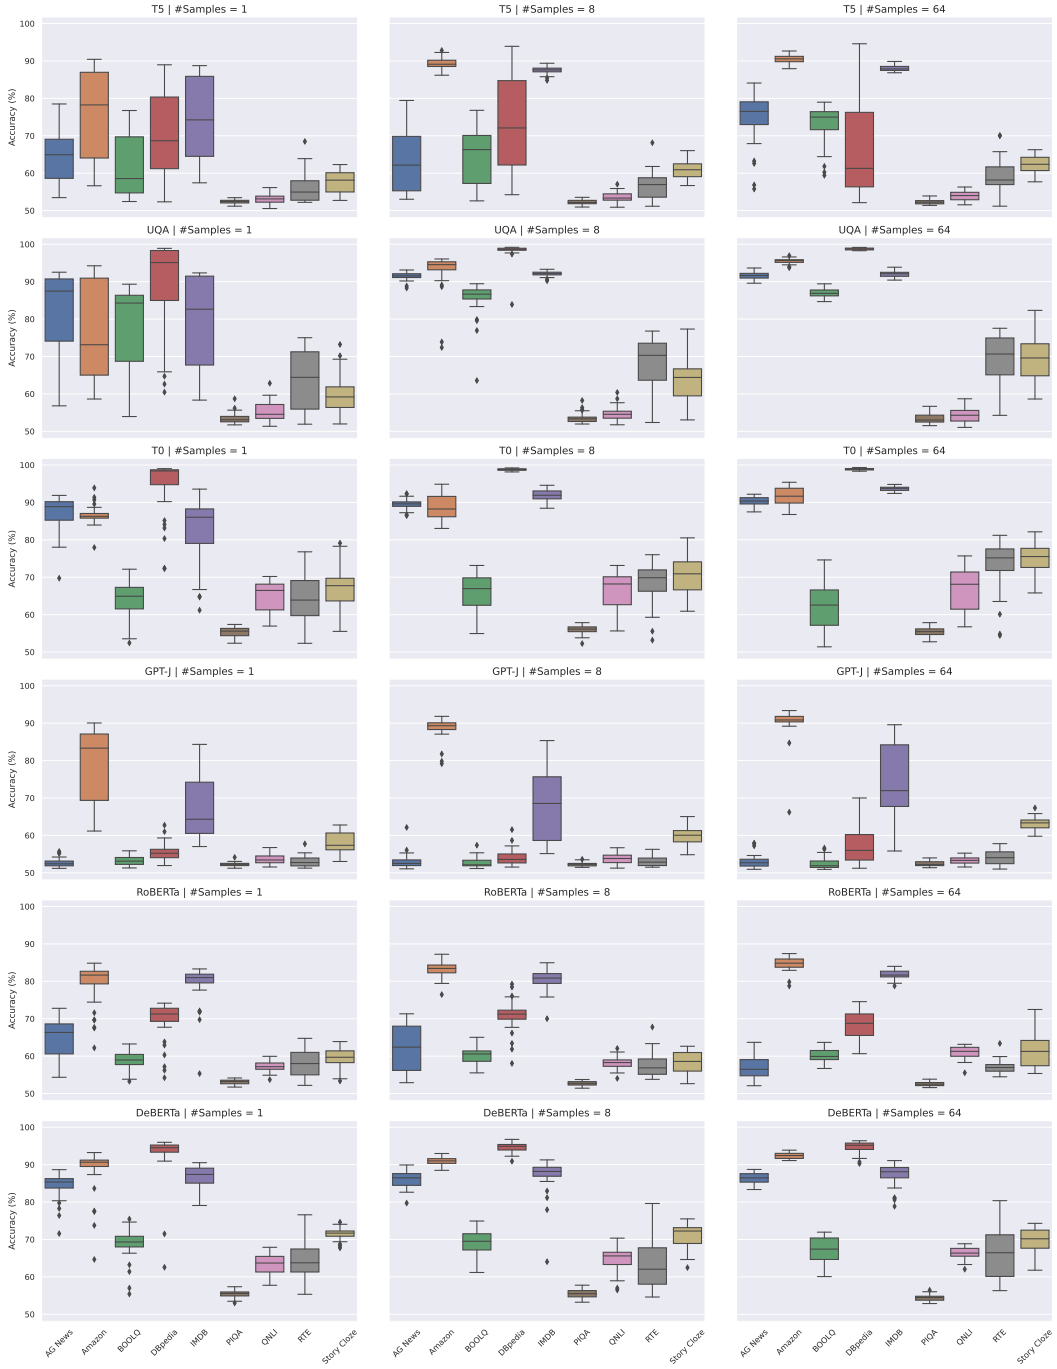

Figure 8: CCS performance when we only use $1, 8, 64$ examples, for all models and datasets. X-axis represents the dataset and y-axis is the accuracy in percentage. For each dataset, we select the prompt with highest 0-shot performance, and then randomly select $1, 8, 64$ data points from this prompt. We then perform CCS and test on all prompts of this dataset, where each value in the barplot corresponds to one prompt.

E Complete Transfer Results

Complete transfer results for CCS, TPC and BSS in all models are shown in Figure 9. Note that we normalize hidden states separately for each dataset when assessing transfer.

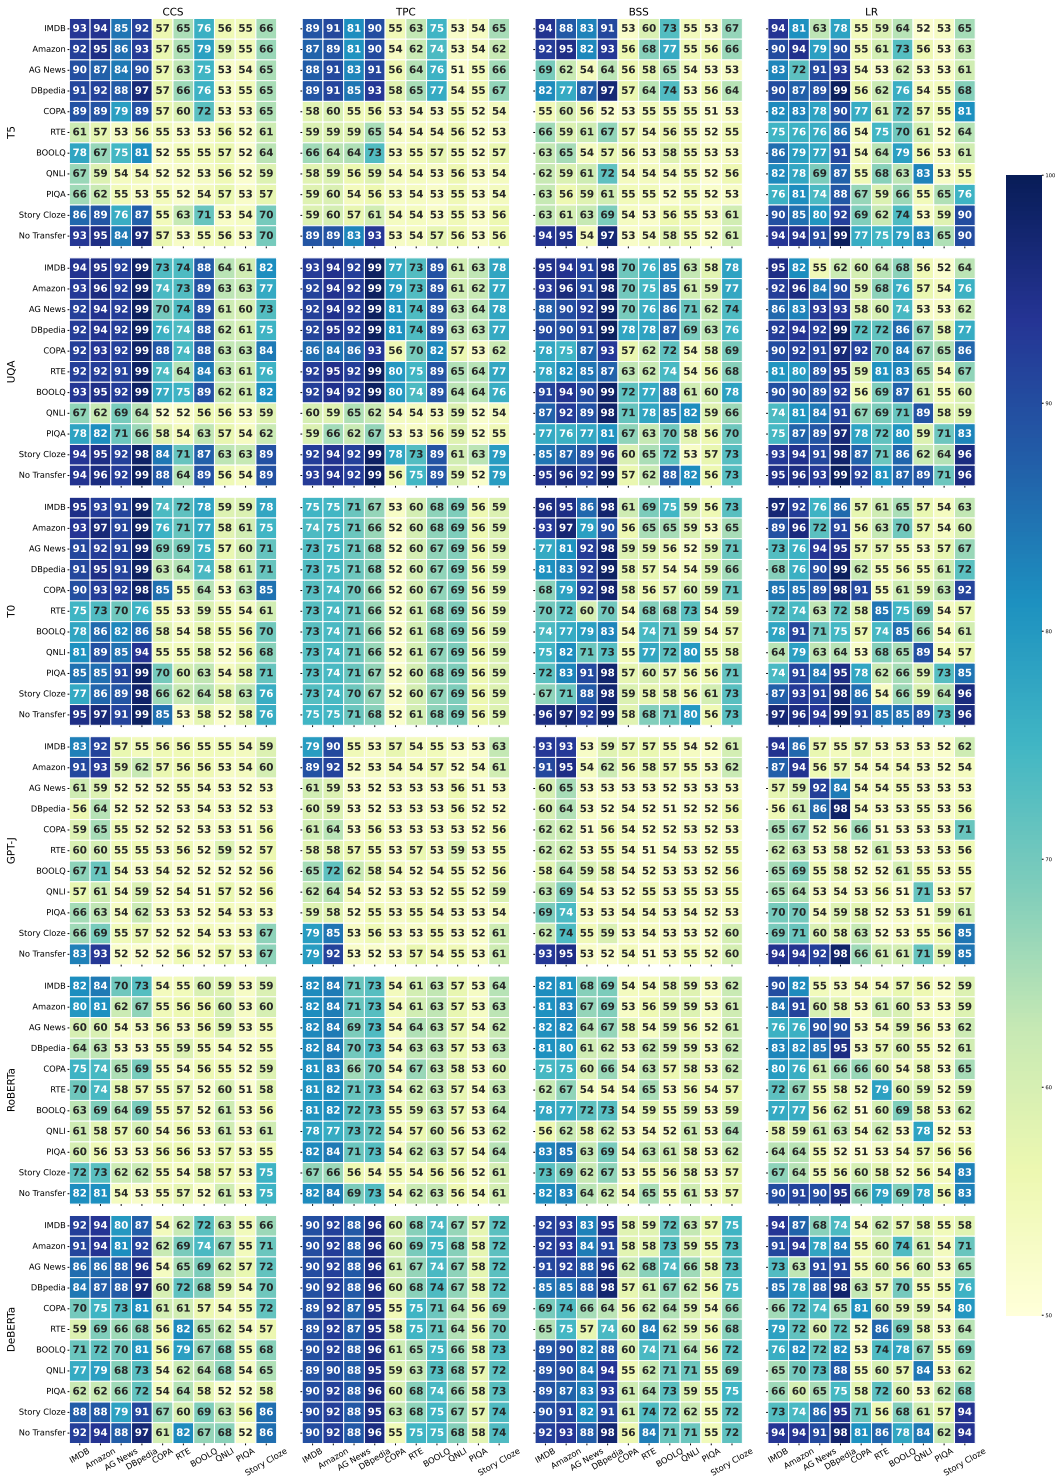

Figure 9: Transfer Accuracy for all models using CCS, BSS, TPC and LR.

F Complete Intermediate Representations Results

In this section we the performance of CCS and LR across all hidden layers in all six models. Specifically, for each model, we generate the hidden states for each dataset every 2 layers (for both encoders and decoders). Then, for each set of hidden states, we perform CCS and LR on each dataset separately, and average the performance across all datasets. We show CCS results in Figure 10 and LR results in Figure 11.

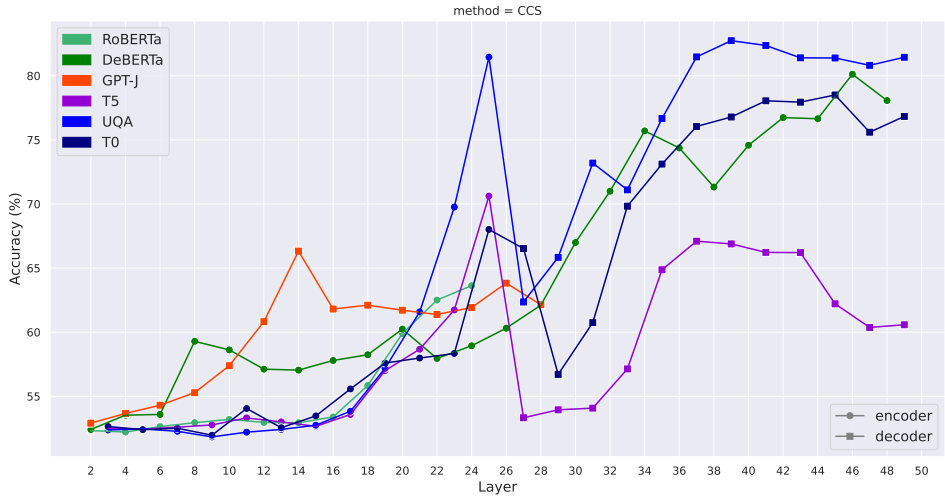

Figure 10: CCS performance when using the hidden states across different layers, using all six models we consider in the paper.

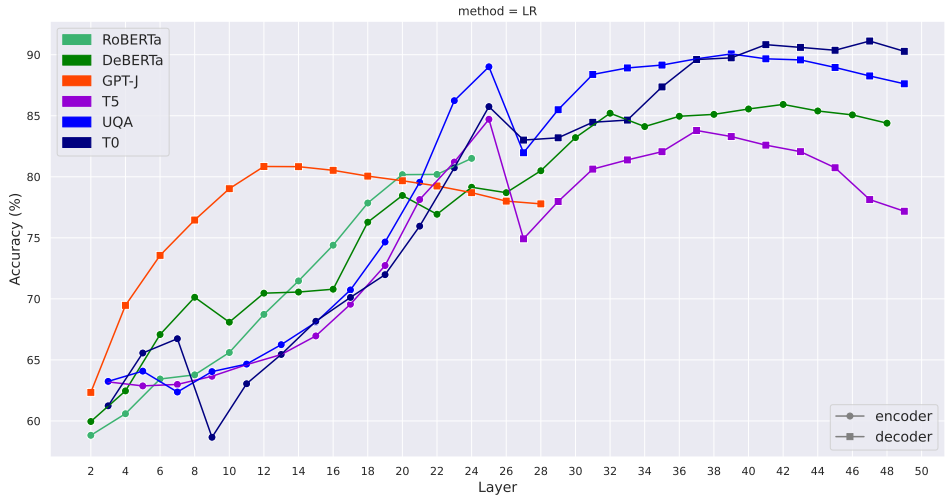

Figure 11: Linear regression performance when using the hidden states across different layers, using all six models we consider in the paper. This is the ceiling of all possible methods, and is supervised.